

Lab 09 — XGBoost: Feature Engineering, Hyperparameter Tuning, and Model Evaluation

Learning goals

- Extract and engineer features from raw datetime and spatial variables
- Apply cyclical (sin/cos) encoding to temporal features
- Fit an XGBoost regression model to predict trip duration
- Understand the regularized XGBoost objective (L1, L2, gamma)
- Tune hyperparameters using randomized search with cross-validation and early stopping
- Evaluate models with regression metrics, diagnostic plots, and partial dependence plots

Lab description

For this lab we use the **Toronto Bike Share** open ridership dataset for **May 2024**. The modelling task is **regression**: predict `trip_duration_min` from temporal and spatial features.

Deliverables

Questions 1–4 answered, notebook and pdf uploaded to Quercus

Setup

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import requests, zipfile, io

from sklearn.model_selection import (train_test_split, cross_val_score,
                                     RandomizedSearchCV)
from sklearn.inspection import PartialDependenceDisplay
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, r2_score
import warnings
warnings.filterwarnings('ignore')
```

Load and prepare data

Run the cells below — these are provided for you.

```
In [2]: base_url = "https://ckan0.cf.opendata.inter.prod-toronto.ca"
package = requests.get(base_url + "/api/3/action/package_show",
    params={"id": "bike-share-toronto-ridership-data"}).json()

for r in package["result"]["resources"]:
    if "2024" in r["name"]:
        url_2024 = r["url"]
        break

resp = requests.get(url_2024)
z = zipfile.ZipFile(io.BytesIO(resp.content))
all_csvs = sorted(f for f in z.namelist() if f.endswith('.csv'))
print("Files in ZIP:", all_csvs)
may_files = all_csvs[4:5]
print("Loading:", may_files)

bike = pd.read_csv(z.open(may_files[0]), encoding='latin-1')
print(f"Shape: {bike.shape}")
bike.head(3)
```

Files in ZIP: ['Bike share ridership 2024-01.csv', 'Bike share ridership 2024-02.csv', 'Bike share ridership 2024-03.csv', 'Bike share ridership 2024-04.csv', 'Bike share ridership 2024-05.csv', 'Bike share ridership 2024-06.csv', 'Bike share ridership 2024-07.csv', 'Bike share ridership 2024-08.csv', 'Bike share ridership 2024-09.csv']
Loading: ['Bike share ridership 2024-05.csv']
Shape: (685782, 11)

```
Out[2]:
```

	Trip Id	Trip Duration	Start Station Id	Start Time	Start Station Name	End Station Id	End Time	End Station Name	
0	28039944	2583	7633	05/01/2024 00:00	Noble St / Queen St W	7548	05/01/2024 00:43	St. Joseph St / Bay St - SMART	3
1	28039945	666	7032	05/01/2024 00:00	Augusta Ave / Dundas St W	7163	05/01/2024 00:11	Yonge St / Wood St	3
2	28039947	1574	7324	05/01/2024 00:00	King St W / Charlotte St	7237	05/01/2024 00:26	Ward St / Wallace Ave	5

```
In [3]: bike.columns = (bike.columns.str.strip()
    .str.replace(r'\s+', '_', regex=True)
    .str.lower())
bike['start_time'] = pd.to_datetime(bike['start_time'])
bike['trip_duration_min'] = bike['trip_duration'] / 60
bike['is_member'] = (bike['user_type'] == 'Annual Member').astype(ir
```

```
bike = bike[(bike['trip_duration_min'] >= 1) & (bike['trip_duration_min'] <= 120)]
print(f"Rows after filtering: {bike.shape[0]:,}")
print(f"Member rate: {bike['is_member'].mean():.1%}")
print("Columns:", list(bike.columns))
```

Rows after filtering: 679,831

Member rate: 12.9%

Columns: ['trip_id', 'trip_duration', 'start_station_id', 'start_time', 'start_station_name', 'end_station_id', 'end_time', 'end_station_name', 'bike_id', 'user_type', 'model', 'trip_duration_min', 'is_member']

```
In [4]: stations = pd.DataFrame(
    requests.get('https://tor.publicbikesystem.net/ube/gbfs/v1/en/station_information.json')['data']['stations'])[['station_id', 'lat', 'lon', 'capacity']]
stations['station_id'] = stations['station_id'].astype(int)
bike = bike.merge(stations, left_on='start_station_id', right_on='station_id')
print(f"Stations matched: {bike['lat'].notna().sum():,} / {len(bike):,}")
```

Stations matched: 658,461 / 679,831

```
In [5]: def haversine_km(lat1, lon1, lat2, lon2):
    R = 6371
    dlat = np.radians(lat2 - lat1); dlon = np.radians(lon2 - lon1)
    a = np.sin(dlat/2)**2 + np.cos(np.radians(lat1))*np.cos(np.radians(lat2))*np.sin(dlon/2)**2
    return R * 2 * np.arcsin(np.sqrt(a))
```

Question 1: Feature Engineering

1a) Extract temporal features

Using `bike['start_time']` (already parsed to datetime), extract the following and add to `bike`:

Column	Description
<code>hour</code>	Hour of day (0–23)
<code>day_of_week</code>	Day of week (0=Mon, 6=Sun)
<code>month</code>	Month (1–12)
<code>is_weekend</code>	1 if Saturday or Sunday, else 0

Print the first 6 rows of these four columns alongside `start_time`.

```
In [6]: # TODO: extract temporal features from start_time
bike['hour'] = bike['start_time'].dt.hour # hint: derive hour-of-day from start_time
bike['day_of_week'] = bike['start_time'].dt.dayofweek # hint: derive weekday from start_time
bike['month'] = bike['start_time'].dt.month # hint: derive calendar month from start_time
bike['is_weekend'] = (bike['day_of_week'] >= 5).astype(int) # hint: build a boolean for weekends

# TODO: print first 6 rows with requested columns
```

```
display_cols = ['start_time', 'hour', 'day_of_week', 'month', 'is_weekend']
print(bike[display_cols].head(6))
```

	start_time	hour	day_of_week	month	is_weekend
0	2024-05-01	0	2	5	0
1	2024-05-01	0	2	5	0
2	2024-05-01	0	2	5	0
3	2024-05-01	0	2	5	0
4	2024-05-01	0	2	5	0
5	2024-05-01	0	2	5	0

1b) Cyclical encoding

Raw integer values like `hour` (0–23) have a discontinuity: hour 23 and hour 0 are adjacent in time but far apart numerically. Apply **sine and cosine** transforms to `hour`, `day_of_week`, and `month` to preserve wrap-around structure.

For a variable with period T :

$$\sin\left(\frac{2\pi x}{T}\right), \quad \cos\left(\frac{2\pi x}{T}\right)$$

Add the six encoded columns to `bike` and print the first 4 rows of `['hour', 'hour_sin', 'hour_cos']`.

```
In [7]: # TODO: cyclical encoding for hour, day_of_week, month
bike['hour_sin'] = np.sin(2 * np.pi * bike['hour'] / 24) # hint: apply sine
bike['hour_cos'] = np.cos(2 * np.pi * bike['hour'] / 24) # hint: apply cosine
bike['dow_sin'] = np.sin(2 * np.pi * bike['day_of_week'] / 7) # hint: apply sine
bike['dow_cos'] = np.cos(2 * np.pi * bike['day_of_week'] / 7) # hint: apply cosine
bike['month_sin'] = np.sin(2 * np.pi * bike['month'] / 12) # hint: apply sine
bike['month_cos'] = np.cos(2 * np.pi * bike['month'] / 12) # hint: apply cosine

# print preview
print(bike[['hour', 'hour_sin', 'hour_cos']].head(4))
```

	hour	hour_sin	hour_cos
0	0	0.0	1.0
1	0	0.0	1.0
2	0	0.0	1.0
3	0	0.0	1.0

1c) Distance to the CN Tower

Using the provided `haversine_km` function, compute the great-circle distance (in km) from each trip's start station to the **CN Tower** (lat = 43.6426, lon = -79.3871). Add the result as `bike['dist_cn_tower_km']` and print the first 4 rows showing `start_station_id`, `lat`, `lon`, and `dist_cn_tower_km`.

```
In [8]: CN_TOWER_LAT, CN_TOWER_LON = 43.6426, -79.3871

# TODO: compute distance from each start station to CN Tower
```

```
bike['dist_cn_tower_km'] = haversine_km(
    bike['lat'], bike['lon'], CN_TOWER_LAT, CN_TOWER_LON) # hint: pass stat

# TODO: print first 4 rows with required columns
print_cols = ['start_station_id', 'lat', 'lon', 'dist_cn_tower_km'] # hint:
print(bike[print_cols].head(4))
```

	start_station_id	lat	lon	dist_cn_tower_km
0	7633	43.642269	-79.429210	3.388666
1	7032	43.652473	-79.401456	1.593569
2	7324	43.645572	-79.393990	0.645427
3	7633	43.642269	-79.429210	3.388666

1d) Assemble the feature matrix

Complete `build_features` to return a DataFrame with:

- `is_weekend`, `is_member` — binary flags
- `hour_sin`, `hour_cos`, `dow_sin`, `dow_cos`, `month_sin`, `month_cos` — cyclical encodings
- `dist_cn_tower_km`, `capacity` — spatial features

Create `X` and `y_reg = bike['trip_duration_min']`. Print the feature count and column names.

```
In [9]: def build_features(df):
    out = pd.DataFrame(index=df.index)

    # TODO: add binary features
    out['is_weekend'] = df['is_weekend'].copy() # hint: copy over the prepared
    out['is_member'] = df['is_member'].copy() # hint: copy over the prepared

    # TODO: add cyclical features
    out['hour_sin'] = df['hour_sin'].copy() # hint: include transformed hour
    out['hour_cos'] = df['hour_cos'].copy() # hint: include transformed hour
    out['dow_sin'] = df['dow_sin'].copy() # hint: include transformed weekday
    out['dow_cos'] = df['dow_cos'].copy() # hint: include transformed weekday
    out['month_sin'] = df['month_sin'].copy() # hint: include transformed month
    out['month_cos'] = df['month_cos'].copy() # hint: include transformed month

    # TODO: add spatial features
    out['dist_cn_tower_km'] = df['dist_cn_tower_km'].copy() # hint: include distance
    out['capacity'] = df['capacity'].copy() # hint: include station capacity

    return out

X = build_features(bike)
y_reg = bike['trip_duration_min'] # hint: use trip duration in minutes as target
print(f'Features: {X.shape[1]}')
print('Columns:', list(X.columns))
```

Features: 10

Columns: ['is_weekend', 'is_member', 'hour_sin', 'hour_cos', 'dow_sin', 'dow_cos', 'month_sin', 'month_cos', 'dist_cn_tower_km', 'capacity']

1e) Data leakage

The raw dataset contains `end_time` and `end_station_id`. Explain in 1–2 sentences why including these as predictors of `trip_duration_min` would constitute **data leakage**, and why the features in 1d) are safe.

Solution:

To avoid data leakage, we should only keep features that could be known at the time of prediction, before observing the outcome. Therefore, we should not include `end_time` and `end_station_id` as predictors of `trip_duration_min` because they are not known at the time of prediction, since `end_time` is `start_time + duration` which directly leaks the outcome and `end_station_id` can only be known after observing the outcome.

Question 2: Baseline XGBoost Model

2a) Train/test split and baseline fit

Split the data into training (70%) and test (30%) sets using `random_state=65`. Fit a baseline `XGBRegressor` with `n_estimators=200`, `max_depth=5`, `learning_rate=0.1`, and `importance_type='gain'`. Report test RMSE and R^2 .

```
In [ ]: # TODO: split data (70/30)
X_train, X_test, y_train, y_test = train_test_split(X, y_reg, test_size=0.3,
print(f"Training rows: {X_train.shape[0]:,} | Test rows: {X_test.shape[0]:,}")

# TODO: define baseline model
xgb_base = XGBRegressor(
    n_estimators=200, # hint: set the number of boosting rounds from the pr
    max_depth=5, # hint: set tree depth from the prompt
    learning_rate=0.1, # hint: set step size from the prompt
    importance_type='gain',
    random_state=65,
    n_jobs=-1
)

# TODO: fit model and predict
xgb_base.fit(X_train, y_train) # hint: train on training features/labels
y_pred_base = xgb_base.predict(X_test) # hint: generate predictions on the

# TODO: compute metrics
rmse_base = np.sqrt(mean_squared_error(y_test, y_pred_base)) # hint: comput
r2_base = r2_score(y_test, y_pred_base) # hint: compute coefficient of dete
print(f'Baseline - RMSE: {rmse_base:.2f} min | R2: {r2_base:.4f}')
```

Training rows: 475,881 | Test rows: 203,950
Baseline – RMSE: 13.54 min | R2: 0.1510

2b) Feature importance

Plot a horizontal bar chart of feature importances (gain) for the baseline model. Which features are most important? Does the ranking make intuitive sense?


Interpretation:

The most important features are `is_member`, `dist_cn_tower_km`, `is_weekend`, `capacity`, and `hour_sin`. This ranking makes intuitive sense because membership status and the distance to the CN Tower can strongly reflect typical trip purpose and behavior, weekends often have different riding patterns than weekdays, and station capacity captures some differences in station sizes. Time of the day matters too, even that is not much but peak or off-peak hours can have different trip durations.

```
In [ ]: # TODO: compute feature importances from fitted baseline model
imp = pd.Series(xgb_base.feature_importances_, index=X.columns) # hint: bui

# TODO: horizontal bar chart
fig, ax = plt.subplots(figsize=(8, 5))
imp.sort_values().plot(kind='barh', ax=ax) # hint: sort importances and plo
ax.set_xlabel('Importance (mean gain)')
ax.set_title('Feature Importance – Predicting trip duration (min)')
plt.tight_layout(); plt.show()

# Optional check
top5 = imp.sort_values(ascending=False).head(5).to_dict() # hint: extract t
print('Top 5 features:', top5)
```



Top 5 features: {'is_member': 0.29598143696784973, 'dist_cn_tower_km': 0.23746848106384277, 'is_weekend': 0.18640518188476562, 'capacity': 0.09490882605314255, 'hour_sin': 0.08027429133653641}

Question 3: Hyperparameter Tuning

3a) Bias-variance tradeoff

Train an `XGBRegressor` with `n_estimators=600`, `max_depth=5`, `learning_rate=0.1`. Using `iteration_range`, compute training and test RMSE at every 10 rounds from 10 to 600. Plot both curves. At roughly what round does test RMSE stop improving? Note this takes about 1 minute to run.

Solution:

Around 300 rounds.

```
In [ ]: # TODO: train model for bias-variance curve
xgb_bv = XGBRegressor(
    n_estimators=600, # hint: use a large round count so you can inspect ov
    max_depth=5, # hint: keep depth fixed as instructed
    learning_rate=0.1, # hint: keep learning rate fixed as instructed
    random_state=65,
    n_jobs=-1
)
xgb_bv.fit(X_train, y_train) # hint: fit once on training data before evalu


n_rounds = list(range(10, 610, 10))
train_err, test_err = [], []

# TODO: compute RMSE at each round using iteration_range=(0, n)
for n in n_rounds:
    tr_pred = xgb_bv.predict(X_train, iteration_range=(0, n)) # hint: predi
    te_pred = xgb_bv.predict(X_test, iteration_range=(0, n)) # hint: predic
    train_rmse_n = np.sqrt(mean_squared_error(y_train, tr_pred)) # hint: ca
    test_rmse_n = np.sqrt(mean_squared_error(y_test, te_pred)) # hint: calc
    train_err.append(train_rmse_n)
    test_err.append(test_rmse_n)

# TODO: plot curves
fig, ax = plt.subplots(figsize=(9, 4))
ax.plot(n_rounds, train_err, 'o-', label='Training RMSE', markersize=3)
ax.plot(n_rounds, test_err, 's-', label='Test RMSE', markersize=3)
ax.set_xlabel('n_estimators (boosting rounds)')
ax.set_ylabel('RMSE (minutes)')
ax.set_title('Bias-Variance Tradeoff: Boosting Rounds vs RMSE')
ax.legend()
plt.tight_layout(); plt.show()

best_n = n_rounds[int(np.argmin(test_err))] # hint: pick the round index wh
```

```
best_test_rmse = min(test_err) # hint: take the minimum value from the test
print(f'Approximate best round: {best_n} (test RMSE: {best_test_rmse:.3f} mi
```



Approximate best round: 600 (test RMSE: 13.468 min)

3b) Randomized search with cross-validation

XGBoost minimizes a regularized objective:

$$\mathcal{L} = \sum_i \ell(y_i, \hat{y}_i) + \sum_k \left[\gamma T_k + \frac{1}{2} \lambda \|w_k\|^2 + \alpha \|w_k\|_1 \right]$$

Use `RandomizedSearchCV` with 5-fold CV on the **training set** to tune `max_depth`, `learning_rate`, `reg_lambda` (L2), and `reg_alpha` (L1). Use `n_iter=20` and `scoring='neg_root_mean_squared_error'`. Print the best parameters and best CV RMSE.

Note: the full grid has $3^4 = 81$ combinations — random search samples 20, reducing compute from 405 to 100 fits.

```
In [15]: param_dist = {
    'max_depth': [4,6,8],
    'learning_rate': [0.05, 0.1, 0.5],
    'reg_lambda': [0.1, 1, 10],
    'reg_alpha': [0, 0.1, 1],
}

# TODO: create base estimator for random search
xgb_rs = XGBRegressor(
    n_estimators=300, # hint: choose a fixed boosting-round budget for each
    random_state=65,
    n_jobs=-1
)

# TODO: set up RandomizedSearchCV (5-fold, n_iter=20, RMSE scoring)
rand_search = RandomizedSearchCV(
    estimator=xgb_rs,
```

```

    param_distributions=param_dist,
    n_iter=20, # hint: number of random parameter samples to evaluate
    cv=5, # hint: number of cross-validation folds
    scoring='neg_root_mean_squared_error', # hint: use negative RMSE so lar
    random_state=65,
    n_jobs=-1,
    verbose=1
)

# TODO: fit on training data only
rand_search.fit(X_train, y_train) # hint: run the randomized search on X_tr

# TODO: print best params and best CV RMSE
best_params_found = rand_search.best_params_ # hint: retrieve selected hype
best_cv_rmse = -rand_search.best_score_ # hint: convert best score back to
print('Best params:', best_params_found)
print(f'Best CV RMSE: {best_cv_rmse:.2f} minutes')

```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

Best params: {'reg_lambda': 10, 'reg_alpha': 0, 'max_depth': 8, 'learning_rate': 0.1}

Best CV RMSE: 13.35 minutes

3c) Early stopping for `n_estimators`

Split the training set into sub-train (80%) / validation (20%). Fit an `XGBRegressor` with `n_estimators=2000`, `max_depth=5`, `learning_rate=0.05`, `early_stopping_rounds=50`. Pass the validation set as `eval_set`. Report the best iteration and validation RMSE.

```

In [ ]: # TODO: split training data into sub-train/validation
X_tr, X_val, y_tr, y_val = train_test_split(X_train, y_train, test_size=0.2,

# TODO: configure model with early stopping
xgb_es = XGBRegressor(
    n_estimators=2000, # hint: set a large upper bound so early stopping ca
    max_depth=5, # hint: use instructed tree depth
    learning_rate=0.05, # hint: use instructed learning rate
    early_stopping_rounds=50, # hint: patience before stopping when validat
    random_state=65,
    n_jobs=-1
)

# TODO: fit with validation eval_set
xgb_es.fit(X_tr, y_tr, eval_set=[(X_val, y_val)], verbose=False) # hint: tr

best_iter = xgb_es.best_iteration # hint: read the best boosting iteration
val_rmse = xgb_es.best_score # hint: read the best validation RMSE reported
print(f'Best round: {best_iter}')
print(f'Val RMSE: {val_rmse:.2f} minutes')

```

Best round: 1367

Val RMSE: 13.30 minutes

Question: Briefly explain why we do early stopping.

Solution:

Early stopping prevents overfitting by terminating the training process when the validation error starts improving minimally, which also tunes `n_estimators` automatically.

3d) Fit the final tuned model

Using the best hyperparameters from 3b), fit a final `XGBRegressor` with `n_estimators=500` on the full training set. Report test RMSE and R^2 . How does it compare to the baseline?

```
In [19]: best_params = rand_search.best_params_ # uncomment if 3b was run
best_params = {'max_depth': 7, 'learning_rate': 0.3, 'reg_lambda': 10, 'reg_

# TODO: fit final tuned model on full training set
xgb_best = XGBRegressor(
    n_estimators=500, # hint: use the instructed round count for final comp
    **best_params,
    random_state=65,
    n_jobs=-1
)
xgb_best.fit(X_train, y_train) # hint: fit tuned model on the full training
y_pred_best = xgb_best.predict(X_test) # hint: produce test-set predictions

# TODO: evaluate tuned model
rmse_tuned = np.sqrt(mean_squared_error(y_test, y_pred_best)) # hint: compu
r2_tuned = r2_score(y_test, y_pred_best) # hint: compute test R-squared for
print(f'Tuned - RMSE: {rmse_tuned:.2f} min | R2: {r2_tuned:.4f}')
print(f'Improvement in RMSE: {rmse_base - rmse_tuned:.2f} min')
```

Tuned - RMSE: 13.45 min | R2: 0.1632
Improvement in RMSE: 0.10 min

Question 4: Model Evaluation

4a) Regression diagnostics

For the tuned model from 3d), create a two-panel figure:

- **Left:** predicted vs actual scatter plot with a 45° reference line
- **Right:** histogram of residuals (actual – predicted) with a vertical line at zero

Comment on any systematic patterns.

Interpretation:

The plots shows over estimation for short trips and under estimation for long trips (predictions shrink toward the mean). The residuals are centered near 0 but right-skewed, meaning some trips are substantially underestimated, so errors are not perfectly symmetric.


```
In [20]: # TODO: regression diagnostics for tuned model
residuals = y_test - y_pred_best # hint: compute actual minus predicted error

fig, axes = plt.subplots(1, 2, figsize=(12, 4))

# Left panel: predicted vs actual
axes[0].scatter(y_test, y_pred_best, alpha=0.3, s=10) # hint: scatter true values
mn, mx = y_test.min(), y_test.max() # hint: use min/max of actual test target
axes[0].plot([mn, mx], [mn, mx], 'r--') # 45-degree line
axes[0].set_xlabel('Actual trip duration (min)')
axes[0].set_ylabel('Predicted (min)')
axes[0].set_title('Predicted vs Actual')

# Right panel: residual histogram
axes[1].hist(residuals, bins=40) # hint: histogram the residual distribution
axes[1].axvline(0, color='red', linestyle='--')
axes[1].set_xlabel('Residual (actual - predicted, min)')
axes[1].set_title('Residual Distribution')

plt.tight_layout(); plt.show()
```



4b) Partial dependence plots

Use `PartialDependenceDisplay.from_estimator` to plot the marginal effect of `is_member`, `dist_cn_tower_km`, and `hour_sin` on predicted trip duration. Create a single figure with three side-by-side panels, one per feature.

Hint: pass `X_train.sample(2000, random_state=65)` instead of the full training set to keep it fast.

Interpret each panel: what does the model say about how each feature relates to trip duration?

Interpretation:


`is_member` shows members are predicted to have shorter trips in general.

`dist_cn_tower_km` generally increases, so farther from the CN Tower leads to longer predicted trips.

`hour_sin` trends slightly downward, which shows that some hours of day are predicted to have slightly shorter trips than others.

```
In [21]: # TODO: PDPs for selected features using a sample for speed
sample_X = X_train.sample(2000, random_state=65) # hint: draw a reproducibl
features_to_plot = ['is_member', 'dist_cn_tower_km', 'hour_sin'] # hint: ir

fig, axes = plt.subplots(1, 3, figsize=(14, 4))
PartialDependenceDisplay.from_estimator(xgb_best, sample_X, features_to_plot)
axes[0].set_title('PDP: is_member')
axes[1].set_title('PDP: dist_cn_tower_km')
axes[2].set_title('PDP: hour_sin')
plt.suptitle('Partial Dependence Plots — Tuned XGBRegressor', y=1.02)
plt.tight_layout(); plt.show()
```



4c) Cross-validation estimate

Use `cross_val_score` with 5-fold CV on the **training set** to estimate generalization RMSE for the tuned model. Compare to the test RMSE from 3d). Are they close?

Interpretation:

The 5-fold CV RMSE on the training set is close to the held-out test RMSE, with a difference of only 0.05 min. This suggests the tuned model's performance is stable and similar across splits in general.

```
In [25]: # TODO: cross-validation RMSE estimate on training set
xgb_cv = XGBRegressor(
    n_estimators=500, # hint: match the tuned model's boosting-round settin
    **best_params,
    random_state=65,
    n_jobs=-1
)

cv_scores = cross_val_score(xgb_cv, X_train, y_train, cv=5, scoring='neg_roc
```

```
cv_rmse = -cv_scores.mean() # hint: convert mean CV score to positive RMSE

print(f'5-fold CV RMSE (training set): {cv_rmse:.2f} min')
print(f'Test RMSE: {rmse_tuned:.2f} min')
print(f'Difference: {abs(cv_rmse - rmse_tuned):.2f} min')
```

```
5-fold CV RMSE (training set): 13.40 min
Test RMSE: 13.45 min
Difference: 0.05 min
```

4d) Summary table

Fill in RMSE and R^2 for both models in the table below.

```
In [26]: # TODO: fill model summary table with your computed metrics
summary = pd.DataFrame({
    'Model': ['Baseline XGBRegressor', 'Tuned XGBRegressor'],
    'Test RMSE (min)': [13.54, 13.45], # hint: plug in rounded baseline/tu
    'Test R2': [0.151, 0.1632], # hint: plug in rounded baseline/t
})
print(summary.to_string(index=False))
```

	Model	Test RMSE (min)	Test R2
Baseline	XGBRegressor	13.54	0.1510
Tuned	XGBRegressor	13.45	0.1632